

Integrity Constraints

- Key constraints:**
- 1. no two distinct tuples of a relation share the same values in all attributes of the key.
 - 2. no proper subset of the key is a superkey.
 - 3. attributes of the key cannot be null.
- Foreign key constraint:** each foreign key value must appear as a primary key value in the referenced relation, or be null.

Relational Algebra


- Unary operators:** selection (σ_c), projection (π_{a_i, a_j}), renaming ($\rho_{a_i:b_i, a_j:b_j}$).
- Binary operators:** set operators ($\cup, \cap, -$), cross-product (\times), (natural) inner join ($R \bowtie_c S$), (natural) outer joins ($R \leftarrow_c S, R \rightarrow_c S, R \leftrightarrow_c S$).
- set operators require relations to be **union compatible**, i.e., have same number of attributes + attributes have same domains (need not have same names).
 - binary operators are left-associative.
 - \times is associative.
 - useful shortcut: $R/S = \pi_A(R) - \pi_A((\pi_A(R) \times S) - R)$.

Entity-Relationship (ER) Model


- Terminology**
- *Entity*: real-world object distinguishable from other objects.
 - *Entity set*: collection of similar entities, represented by **rectangles**.
 - *Attribute*: specific information describing an entity, represented by **ovals**.
 - *Key*: represented as **underlined** attributes.
 - *Relationship*: association between two or more entities.
 - *Relationship set*: collection of similar relationships, represented by **diamonds**.
 - *Relationship roles*: shown explicitly when one entity set appears ≥ 2 times in a relationship set.
 - *Relationship keys*: the primary key of a relationship set consists of the keys of its entities, as well as its underlined attributes.
 - *Degree*: a relationship set with degree n (i.e., n -ary relationship set) involves n entity roles.
 - *Aggregation*: can be used to model relationships between an entity set and a relationship set by representing relationships as higher level entity sets. This is indicated using a **bold square**.

ER constraints


- *Key constraint*: each instance of E can participate in at most one instance of R .



- *Total participation constraint*: each instance of E must participate in at least one instance of R .



- *Partial participation constraint*: each instance of E can participate in 0 or more instances of R .

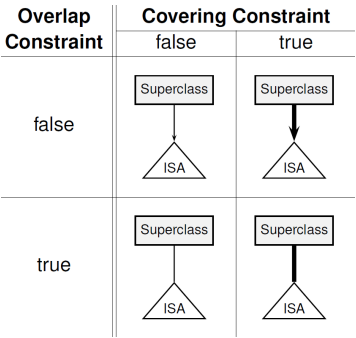


- *Weak entity set*: E does not have its own key, and requires the primary key of its *owner entity* E' to be uniquely identified. There must be a many-to-one (identifying) relationship from E to E' , and a total participation on R .



ISA hierarchies

- We can classify entity sets into subclasses. Every entity in a subclass entity set is an entity in its superclass entity set. ISA hierarchy constraints include:
- *Overlap constraint*: can an entity belong to multiple subclasses?
 - *Covering constraint*: should an entity belong to some subclass?



SQL

Data types: boolean, integer, numeric/float, char(n), varchar(n), text, date, timestamp.

Basic syntax

- **Create table:**
create table [if not exists] T (id integer primary key);
- **Drop table:**
drop table if exists T [cascade];
- **Modify table:**
insert into T values (X, Y);
delete from T [where id = X];
update T set id = id * X [where name = Y];
alter table T add/drop/alter column/constraint X [...];
 - if a value is not specified for a column when a tuple is added, the default value for the column would be assigned.
- **Conceptual evaluation of queries:**
select [distinct] select-list
from from-list
where where-condition
group by group-by-list
having having-condition
order by order-by-list
offset offset-spec [asc (default)/desc]
limit limit-spec
 1. Compute cross-product of tables in from-list.
 2. Select tuples evaluating to *true* for the where-condition.
 3. Partition selected tuples using the group-by-list.
 4. Select groups evaluating to *true* for the having-condition.
 5. Generate an output tuple for each selected group based on the select-list.
 6. Remove duplicate output tuples (if distinct is specified).
 7. Sort output tuples based on the order-by-list.
 8. Remove output tuples based on the offset-spec and limit-spec.

Validity of group by/having (one of the following must hold):

 - Output column A appears in the group by clause.
 - A appears in an aggregated expression in the select/having clause.
 - the primary key of relation R appears in the group by clause.
- **Transactions:**
begin;
[set constraints CONSTR_NAME/TRIG.NAME deferred;]
...sql statements...
commit/rollback;
Transactions adhere to the ACID properties:
 - *Atomicity*: either all effects are reflected or none.
 - *Consistency*: executed in isolation, preserves integrity constraints of database.
 - *Isolation*: no interference with other concurrent transactions.
 - *Durability*: effects persists even if system fails.
- **CTEs:**
with CTE1 as (...),
CTE2 as (...),
CTE3 as (...)
select ...;
- **Views:**
create or replace view V(A, B) as ...;

- **Conditional expressions:**
case
 when X = Y then A
 else B
end as VAR

case X
 when Y then A
 else B
end as VAR

Other syntax

- **Schema constraints:**
[constraint CONSTR_NAME] CONSTRAINT [deferrable initially deferred/immediate]
 - possible CONSTRAINTS:
 - * unique
 - * not null
 - * primary key / unique not null
 - * foreign key (A, ...) references T [match full] [on update/delete ACTION]
 - * check ...
 - * default '...'
 - match full: prevents reference to completely or partially null entries.
 - possible ACTIONS:
 - * no action: rejects delete/update if it violates constraint.
 - * restrict: same as above, but constraint checking cannot be deferred.
 - * cascade: propagates delete/update to referencing tuples.
 - * set default: updates FKs of referencing tuples to some default value.
 - * set null: updates FKs of referencing tuples to null.
- **Set operations:**
union/intersect/except [all]
 - all preserves duplicate records.
- **Join operations:**
[natural] [left/right/full] join
- **Comparison predicates:**
is [not] [null/distinct from]
- **Subquery expressions:**
exists, in, any/some, all, unique
- **Aggregate functions:**
count, sum, avg, min, max
- **Coalesce:** returns the first non-null value.
select coalesce(a1, a2, ...) as A;
- **Nullif:** returns null if a1 = a2, else returns a1.
select nullif(a1, a2) as A;
- **Pattern matching:**
... where A like '%';
 - .: matches any single character.
 - %: matches any sequence of 0 or more characters.
- **Miscellaneous:**
 - 'a' || 'b': string concatenation.
 - round(A): rounding off to nearest integer value.

Functions

- create or replace function F(<param> <type>, ...) returns <type> as \$\$
...
\$\$ language sql;
select F(...); / select * from F(...);
- **Return single tuple:**
function F(...) returns T as ...
 - **Return single custom tuple:**
function F(IN ..., OUT ...) returns <name> as ...
 - **Return ≥ 2 tuples:**
function F(...) returns **setof** T as ...
function F(...) returns **table**(<param> <type>, ...) as ...
 - **Return nothing:**
function F(...) returns **void** as ...

Procedures

```
create or replace procedure P(<param> <type>, ...)
as $$ ... $$ language sql;
call P(...);
```

PL/pgSQL

Basic syntax

```
• Variables:
declare
    varA integer := 0;
    varB integer;
begin
    select id into varC from T;
    ...
end;
$$ language plpgsql;

• Conditional statements:
if c1 then ...;
elsif c2 then ...;
else ...;
end if;

• Cursors:
declare
    curs cursor for (select * from T where ...);
    var record;
begin
    open curs;
    loop
        fetch curs into var;
        -- assign output column names to values
        return next;
    end loop;
    close curs;
end;
fetch [prior/first/last/absolute n/relative n] from curs into var;
move [prior/first/last/absolute n/relative n] from curs;
[update/delete] <table> ... where current of curs;
```

Triggers

```
create or replace function TRIG.F()
returns trigger as $$
begin ... end;
$$ language plpgsql;
---
create trigger TRIG.NAME
[after/before/instead of] [insert/delete/update] on T
for each [row/statement]
[when <condition>] -- trigger condition
execute function TRIG.F();
• instead of triggers can only be defined on views and on the row level.
```

Special values

- NEW: new tuple being inserted or updated. null for deletion.
- OLD: old tuple being deleted or updated. null for insertion.
- CURRENT_DATE: current date.
- TG.OP: operation that activates the trigger.
- TG.TABLE_NAME: name of table causing trigger invocation.

Return values

- before insert: null = no insertion; else output inserted.
- before update: null = no update; else tuple updated with output.
- before delete: null = no deletion; else deletion proceeds as normal.
- after <action>: return value does not matter.
- instead of <action>: null = ignore operations on current row; else proceed as normal.
- statement-level triggers: return values ignored, use raise exception ‘...’ to omit subsequent operations.

Trigger condition requirements

- No select in when().
- No OLD in when() for insert.
- No NEW in when() for delete.
- No when() for *instead of* triggers.

Deferred triggers

```
create constraint trigger TRIG.NAME
after [insert/delete/update] on T
deferrable initially [deferred/immediate]
for each row
execute function TRIG.F();
---
begin transaction;
set constraints TRIG.NAME deferred;
update ...;
commit;
• Deferred triggers only work with after and for each row.
```

Order of trigger activation

- before statement-level triggers.
- before row-level triggers.
- after row-level triggers.
- after statement-level triggers.

Within each category, triggers are activated in *alphabetical order*. Also note that if a before *row-level* trigger returns null, all subsequent triggers on the same row will be omitted.

Normal Forms

Normal forms help to reduce redundancy in relational schemas.

Functional dependencies (FDs)

$A \rightarrow B$: if two rows have the same A , then they have the same B .

- Reflexivity:** $AB \rightarrow A$.
- Augmentation:** if $A \rightarrow B$, then $AC \rightarrow BC$.
- Transitivity:** if $A \rightarrow B$ and $B \rightarrow C$, then $A \rightarrow C$.
- Decomposition:** if $A \rightarrow BC$, then $A \rightarrow B$ and $A \rightarrow C$.
- Union:** if $A \rightarrow B$ and $A \rightarrow C$, then $A \rightarrow BC$.

Closures

$\{A\}^+$: denotes the set of attributes that can be decided by A .
To prove $A \rightarrow B$, it suffices to show that $B \in \{A\}^+$.

Algorithm for computing closures

- Initialize the closure to $\{A\}$.
- If there is a FD such that $A \rightarrow B$, we put B into the closure.
- Repeat step 2 until no new attributes can be added.

Keys

- Superkey:** set of attributes that decides all other attributes.
- Key:** minimal superkey.
- Prime attribute:** attribute which appears in a key.

Algorithm for deriving keys of T

- Consider all attribute subsets in T .
- Compute the closure of each subset.
- Identify all superkeys based on the closures.
- Identify all keys.

Useful heuristics:

- Check for attribute sets from smallest to largest.
- If an attribute does not appear in the R.H.S. of any FD, it must be in every key.

Non-trivial and decomposed FDs

- Decomposed FD:** an FD whose R.H.S. has only one attribute.
- Non-trivial FD:** an FD whose attributes in the R.H.S. do not appear in the L.H.S.

Algorithm for deriving non-trivial and decomposed FDs of T

- Consider all attribute subsets in T .
- Compute the closure of each subset.
- Remove all trivial attributes from each closure.
- Decompose all FDs (using rule of decomposition).

Boyce-Codd Normal Form (BCNF)

T is in BCNF if every non-trivial and decomposed FD has a superkey as its L.H.S. A table violates BCNF if it does not satisfy the “*more but not all*” condition, i.e., there exists a closure of size k which:

- contains more than k attributes, but
- does not contain all the attributes in T .

BCNF properties

- (+) no update, deletion, or insertion anomalies.
- (+) few redundancies.
- (+) original table can be reconstructed from decomposed tables.
- (+) guarantees **lossless join decomposition** (i.e., common attributes in decomposed tables always constitute a superkey of either sub-table).
- (−) does not guarantee **dependency preservation** (i.e., set of FDs on the original table and set of FDs on the decomposed tables may not be equivalent).

Algorithm for BCNF decomposition

- Find a subset X of attributes in T whose closure $\{X\}^+$ satisfies the “more but not all” condition.
- Decompose T into two tables T_1 and T_2 such that:
 - T_1 contains all attributes in $\{X\}^+$, and
 - T_2 contains all attributes in X and all attributes not in $\{X\}^+$.
- Check if T_1 and T_2 are in BCNF.
- If needed, recursively decompose either T_1 or T_2 .

Third Normal Form (3NF)

T is in 3NF if every non-trivial and decomposed FD either:

- has a superkey as its L.H.S., or
- has a prime attribute as its R.H.S.

Minimal basis (MB)

A *minimal basis* of a set S of FDs is a simplified version of S satisfying **all** following conditions:

- Every FD in MB can be derived from S and vice versa.
- Every FD in MB is a non-trivial and decomposed FD.
- No FD in MB is **redundant** (i.e., can be derived from other FDs in the MB).
- Every FD in MB does not have **redundant attributes** in its L.H.S. (i.e., if an attribute is removed from the L.H.S., then the FD cannot be derived from S).

Algorithm for minimal basis

- Decompose all FDs (using rule of decomposition).
- Remove redundant attributes on the L.H.S. of each FD.
- Remove redundant FDs.

Algorithm for 3NF decomposition

- Derive a MB for the set of FDs on T .
- Combine FDs whose L.H.S. are the same in the MB.
- Create a table for each remaining FD.
- If none of the tables contain a key of T , create a table that contains a key of T .

Algorithm for verifying BCNF/3NF satisfiability

- Compute the closure of all attribute subsets in T .
- Derive the keys of T .
- For each closure, check if:
 - it satisfies the “more but not all” condition, and
 - its additional attribute(s) is not a prime attribute. (3NF only)
- If such a closure exists, T is not in BCNF/3NF.

Note: Alternatively, for step 3(a), we could check if the L.H.S. of each FD of T is a superkey.