

CS4226 Internet Architecture

Probability

- **Sample space**,  $S$ : the set of all outcomes.
- **Event**,  $E$ : a subset of the sample space.
- **Probability function**,  $P(E)$ :
  1.  $0 \leq P(E) \leq 1$ .
  2.  $P(S) = 1$ .
  3.  $P(\bigcup_{i=1}^\infty E_i) = \sum_{i=1}^\infty P(E_i)$  for any **mutually exclusive** events  $E_1, E_2, \dots$ .
- **Conditional probability**:  $P\{E_i|E_j\} := \frac{P\{E_i \cap E_j\}}{P\{E_j\}}$ .
  - **Joint probability**:  $P\{E_i \cap E_j\} = P\{E_i|E_j\}P\{E_j\}$ .
  - **Bayes' rule**:  $P\{E_i|E_j\} = \frac{P\{E_j|E_i\}P\{E_i\}}{P\{E_j\}}$ .
- **Law of total probability**: let events  $E_1, \dots, E_n$  be a *partition* of the sample space, i.e.,  $\bigcup_i E_i = S$  and  $\forall i, j, E_i \cap E_j = \emptyset$ . For any event  $E$ ,

P{E} = P{E ∩ S} = P {⋃\_i (E ∩ E\_i)} = ∑\_i P {E ∩ E\_i} = ∑\_i P {E|E\_i} P {E\_i}.

- **Random variable**,  $X$ : a function that assigns a real value to each outcome  $s \in S$ .
  - For any set of real numbers  $A \subseteq \mathcal{R}$ ,  $P\{X \in A\} := P(X^{-1}(A))$ .
  - **Distribution function**:  $F(x) := P\{X \leq x\} = P(X^{-1}((-\infty, x]))$ .
  - An r.v. is continuous if there exists a **probability density function** s.t.  $f(x) := \frac{d}{dx} F(x)$ .
  - **Independence**: 2 rvs are independent if the realization of one does not affect the probability distribution of the other,  $f_{X,Y}(x, y) = f_X(x)f_Y(y)$ .
- **Expectation/Mean** of a rv  $X$ :  $E[X] := \begin{cases} \int_{-\infty}^\infty xf(x)dx & \text{(continuous rv)} \\ \sum_{x=-\infty}^\infty xP\{X = x\} & \text{(discrete rv)} \end{cases}$
- **Exponential distribution**: a continuous rv  $T$  follows an exp. distribution with parameter  $\lambda > 0$  if for  $x \geq 0$ ,  
 $F(x) = P\{T \leq x\} = 1 - e^{-\lambda x}$  (or  $F(x) = P\{T > x\} = e^{-\lambda x}$ )  
 $f(x) = \frac{dF(x)}{dx} = \lambda e^{-\lambda x}$   
 $E[T] = \int_{-\infty}^\infty xf(x)dx = \frac{1}{\lambda}$ 
  1. **Memoryless property**:  $P\{T > s + t | T > s\} = P\{T > t\}$ .
  2. **Merging property**: merging 2 arrival patterns gives rise to a process which is also Poisson, with  $T = \min(T_1, T_2)$  and rate  $\lambda = \lambda_1 + \lambda_2$ .
  - Comparison between rvs:  $P\{X > Y\} = \frac{\lambda_Y}{\lambda_X + \lambda_Y}$ .
- **Geometric distribution**: discrete distribution with density function  $P\{L = i\} = p(1 - p)^{i-1}$ .
  - Used to model the random number of coin flips of tails until you get a head.
  - 1. **Memoryless property**:  $P\{L > s + t | L > s\} = P\{L > t\}$ .


Performance

- **Link rate/capacity/bandwidth**: # bits that can be pushed onto a link per unit time.
- **Throughput**: # bits that can be communicated (i.e., transferred e2e) per unit time.
- **End-to-end delay**: related to
  - a. Packet size.
  - b. Queue size.
  - c. Bandwidth.
  - d. Distance.
  - e. Propagation speed.Comprises of:
  1. **Processing delay**: time for processing data into packets, affected by [a].
  2. **Queueing delay**: affected by [b].
  3. **Transmission delay**: time taken to place data onto the link, affected by [a,c].
  4. **Propagation delay**: time taken by the last bit of the packet to reach the destination, affected by [d,e].
- **Little's law**:  $L = \lambda W$ .
  - **Average # customers in system**:  $L = \lim_{t \rightarrow \infty} \frac{1}{t} \int_0^t L(s) ds$  (as a deterministic limit).
    - \*  $L(t)$ : # customers in the system at time  $t$ .
  - **Arrival rate**:  $\lambda = \lim_{t \rightarrow \infty} \frac{N(t)}{t}$  (as a deterministic limit).
    - \*  $N(t)$ : # customers arrived till time  $t$ .
  - **Average sojourn time**:  $W = \lim_{n \rightarrow \infty} \frac{1}{n} \sum_{j=1}^n W_j$  (as a deterministic limit).

Network Queueing Models

- **M/M/1 model**: M: Markovian arrival time, M: Markovian service time, 1: single server.
  - Assumes a **queue of infinite size** (i.e., no packets are dropped).
  - **Arrival time**: Poisson with rate  $\lambda$ .
  - **Service time**: exponential i.i.d. with rate  $\mu$ .
  - Arrival and service times are **independent**.
  - **FIFO** service discipline.
- Notations:
  1. Arrival:
    - **Arrival time** of  $i^{\text{th}}$  packet:  $t_i$ .
    - **Inter-arrival time**:  $T_i = t_{i+1} - t_i$ .
      - \*  $T_i$ 's are **independent and identically distributed (i.i.d.) exponential** as  $T$  with mean  $1/\lambda$ .
      - \* The arrival pattern is a **Poisson process**.
    - **Arrival rate**:  $\lambda = 1/E[T]$ .
  2. Service:
    - **Service time**,  $S_i$ : follows i.i.d. rv  $S$ , with mean  $1/\mu$ .
    - **Service/Processing rate**:  $\mu = 1/E[S]$ .


- Results:
  1. **Utilization**,  $\rho$ : % time server is busy, or probability that a random observation finds the server busy,  
 $\rho = \frac{\lambda}{\mu}$ .
    - For system stability, we require  $\lambda < \mu$ .
    - If  $\lambda \geq \mu$ , the system will be permanently utilized, and the queue will grow indefinitely.
  2. **Percentage of time that exactly  $i$  packets are in the system**,  $\pi_i$ :  
 $\pi_i = P\{L = i\} = \rho^i (1 - \rho)$ .
    - $\pi = (\pi_0, \pi_1, \dots)$ .
    - $\sum \pi_i = \frac{1-\rho}{1-\rho} = 1$  (∴ sum of infinite geometric series,  $S = \frac{a}{1-r}$ ).
    - $\pi_i$  follows a **geometric distribution**.
  - Relationship to Little's law ( $L = \lambda W$ ):
    1. **Average # packets in an M/M/1 queueing system**:  $E[L] = E[Q] + E[X]$  (as an rv).
      - $E[Q]$ : avg # packets in queue.
      - $E[X]$ : avg # packets in server.
        - \*  $E[X] = \sum_{x=0}^\infty xP\{X = x\} = P\{X = 1\}$ , since the M/M/1 model assumes a single server.  
 $P\{X = 1\} = P\{L - Q = 1\}$  # in server = # in system - # in queue  
 $= P\{L > 0\}$  see derivation below  
 $= 1 - \pi_0$  ∴  $\pi_0$  rep. % time system is idle  
 $= \rho$  ∴  $\pi_0 = 1 - \rho$   
Thus,  $E[X] = \rho$  also rep. the probability of the system being busy.
        - \* Derivation of  $\{L - Q = 1\} = \{L > 0\}$ :  
 $\{L - Q = 1\} \subseteq \{L > 0\}$  ∴  $L = 1 + Q \Rightarrow L > 0$   
 $\{L > 0\} \subseteq \{L - Q = 1\}$  ∴ server cannot be idle when queue is nonempty  
∴ if system is nonempty, some pkt is in the server
    2. **Arrival rate**:  $\lambda = 1/E[T]$  (as a function of an rv).
    3. **Sojourn time**:  $W_i = t_i^d - t_i$ .
      - $t_i^d$ : **departure time** of  $i^{\text{th}}$  packet.
      - In a queueing system,  $W_i = D_i + S_i$  (as an rv).
        - \*  $D_i$ : sojourn time in queue (i.e.,  $E[D]$  = mean queueing delay).
        - \*  $S_i$ : sojourn time in server (i.e.,  $E[S]$  = mean service time).
  - From the perspective of the system (i.e., queue + server), we have:
    - **Average # packets in system**:  $E[L] = \frac{\rho}{1-\rho}$ .
    - **Average sojourn time** for single entrance:  $E[W] = \frac{1}{\mu - \lambda}$ .
      - \* **Higher capacity**  $\mu$  accommodates **higher throughput**  $\lambda$  and provides **lower delay**  $E[W]$ .



- As  $\rho = \lambda/\mu \rightarrow 1, \lambda = \mu \Rightarrow$  waiting time tends toward  $\infty \Rightarrow$  unstable system.
- From the perspective of the queue, we have:
    - **Average # packets in queue**:  $E[Q] = \frac{\rho^2}{1-\rho}$ .
    - **Average queueing delay**:  $E[D] = E[W] - \frac{1}{\mu}$ .

Statistical multiplexing vs TDM

- **Bandwidth**:
  - **Physical link capacity**: the theoretical processing limit of the hardware.
  - **Effective bandwidth**: actual achievable throughput on the link, depends on utilization  $\rho$  and quality of service.
- **Packet-switching networks**:
  1. **Store-and-forward**: entire packet must arrive at the router before it can be forwarded.
  2. **Statistical multiplexing**: link bandwidth is shared on demand.
    - Distinct from **time-division multiplexing (TDM)** used in circuit-switching networks, where each host gets the same slot in a revolving TDM frame.




- (+) **More users** are allowed to use the network, as compared to circuit-switching.
- (+) Great for (asynchronous) bursty data.
- (-) **Excessive congestion**: protocols are needed for reliable data transfer and congestion control.
- **Circuit-switching networks**:
  - In **TDM**, each Poisson stream has its own queue, with params  $(\frac{\lambda}{k}, \frac{\mu}{k})$ . Compared to statistical multiplexing,
    1. Both have the **same server utilization**:  $\rho = \frac{\lambda}{\mu} = \frac{\lambda/k}{\mu/k}$ .
    2. **Average sojourn time** is  $k$  times larger under TDM, since  $E[W] = \frac{1}{\mu - \lambda} = \frac{1}{\mu} \frac{1}{1-\rho}$ .
    3. **Average # queueing pkts** is  $k$  times larger under TDM, since  $E[Q]$  is the same with SM for each sub-queue.
    4. **Total # pkts in system** is  $k$  times larger under TDM, since  $E[L]$  is the same with SM for each sub-queue.  
∴ **SM is more efficient than TDM** in terms of resource utilization: (1) lower sojourn time, (2) lower average # packets in system and queue.

Tandem queues


- Burke's theorem**: if an M/M/1 system with arrival rate  $\lambda$  **starts in a steady state**,
1. the **departure process** is Poisson with rate  $\lambda$ , and
  2. **# of customers in the system** at any time  $t$  is *independent* of the sequence of departure times prior to  $t$ .
- When queues are in tandem:
    - **Utilization** of each server  $i$ :  $\rho_i = \frac{\lambda}{\mu_i}$ .
    - **Joint probability**  $P\{\dots, L_k = l_k, \dots\} = \prod_i P\{L_i = l_i\} = \prod_i \rho_i^{l_i} (1 - \rho_i)$  by independence.
      - \* **Average # jobs in system**:  $E[L] = \sum_i E[L_i]$ .
  - With probabilistic routing, different flows might take different paths.
    - **Utilization** of each server  $i$ :  $\rho_i = \frac{\lambda_i}{\mu_i}$ .
      - \*  $\mu_i$  is known; the service rate of each subsystem  $i$ .
      - \*  $\lambda_i$ , the **effective arrival rate**, is given by:  
 $\lambda_i = r_i + \sum_{j=1}^n \lambda_j P_{ji}$  (equivalently,  $\lambda = r(I - P)^{-1}$  in matrix form)
        - $r_i$ : rate of external arrival to server  $i$ .
        - $\lambda_j$ : output rate of subsystem  $j$  (= effective arrival rate of  $j$ , assuming the system is in steady state).
        - $P_{ji}$ : fraction of packets arriving from subsystem  $j$  to subsystem  $i$ .

- **Jackson network**:



Resource Allocation

- The Internet provides **best-effort service** (i.e., no guarantee for packet delivery).
  - Suitable for **elastic traffic**, but not for **real-time applications**.
  - Solution: allocate *different amount of resources* to different application flows.
    - \*  $U_i$ : utility function based on the type of application that is running.



- Notations:
  - $C = (C_1, \dots, C_n)$ : **resource capacity** of each link.
  - $d = (d_1, \dots, d_m)$ : **flow demand**.
  - $x = (x_1, \dots, x_m)$ : **feasible solution** satisfying:
    1.  $\forall i, 0 \leq x_i \leq d_i$ , and
    2. each flow  $j$  going through link  $i$   $x_j \leq C_i$ .
  - $\phi = (\phi_1, \dots, \phi_m)$ : **weight vector** containing weights for each flow.
- **Fairness**:
  1. Users with small demand get all they want (and not more).
  2. Users with large demand evenly split the remaining resources.
- **Max-min fairness**: a feasible allocation is max-min fair iff an *increase of any rate* within the feasible domain results in a *decrease of a smaller or equal rate*.
  - i.e.,  $x$  is max-min fair iff for any feasible  $y$ , if  $y_i > x_i$ , then  $\exists j$  s.t.  $y_j < x_j \leq x_i$ .
  - There always exists a *unique* max-min fair solution.
  - **Bottleneck**: we say that resource  $r$  is a bottleneck for flow  $i$  iff:
    1. Resource  $r$  is saturated (i.e., not possible to increase any  $x_i$ ), and
    2. Flow  $i$  has the maximum rate  $x_i$  among all flows using resource  $r$ .
  - **Theorem**: when each flow has an *infinite demand* under a network system, a flow allocation is max-min fair iff *every flow has a bottleneck resource*.
  - **Water filling algorithm**:
    1. List the demands as empty buckets with width 1 and area  $d_i$ .
    2. Find the fair share at each bottleneck resource, and sort in ascending order.
    3. Starting from the lowest fair share, fill in water until a bottleneck is reached.
    4. For the remaining flows, repeat step 3 for each bottleneck resource.

### Weighted max-min fair share:

- Users receive a fair share based on their weights,  $x_i = \frac{\phi_i}{\sum_j \phi_j} C$ .
- Spillover resources (i.e.,  $\sum_i x_i - d_i$  for each  $i$  with  $x_i > d_i$ ) are distributed to the remaining flows, proportional to the weights of each flow.
- Repeat step 2 until no spillover resources are left.
- Bottleneck:** we say that resource  $r$  is a bottleneck for flow  $i$  iff:
  - Resource  $r$  is saturated (i.e., not possible to increase any  $x_i$ ), and
  - Flow  $i$  has the max *normalized* rate  $\frac{x_i}{\phi_i}$  among all flows using resource  $r$ .
- Water filling algorithm** for weighted max-min fair share:
  - List the demands as empty buckets with **width**  $\phi_i$  and **area**  $d_i$ .
  - Find the fair share at each bottleneck resource, and sort in ascending order.
  - Starting from the lowest fair share, fill in water until a bottleneck is reached.
  - For the remaining flows, repeat step 3 for each bottleneck resource.

## Software Defined Networking (SDN)

- Approach to networking (i.e., **network implementation + management**) based on new fundamental principles.
  - Note: traditional networking is *protocol-defined*, not *software defined*.
- Motivations:**
  - While the Internet guarantees *best-effort* packet delivery and enables innovation in applications, change is only easy at the **edge** but not at the **core**.
    - Network edge:** an hourglass IP model providing **layered service abstraction**.
      - Decomposes delivery into fundamental components, enabling independent innovation at each layer.
    - Network core:** limitations include
      - Closed equipment:** software are bundled with hardware, and interfaces are vendor-specific.
      - Over-specification:** contributes to slow protocol standardization.
      - Restricted innovation:** only vendors can improve the system  $\rightarrow$  long delays in introducing new features.
      - Difficulty in network management:** high cost of network operation + likelihood of buggy equipment.
  - As compared to **computation and storage** (which have been virtualized), networks are lagging behind.
    - Management is complex and primitive**, and rely heavily on network administrators.
    - Routing algorithms change very slowly**.
    - New control requirements (e.g., isolation, traffic engineering, pkt processing + analysis, etc.)  $\rightarrow$  **complexity**.
- Routers** have 2 key functions:
  - Routing:** run routing algorithms/protocols (e.g., RIP, OSPF, BGP).
    - The **control plane** establishes the route state, and determines how/where packets are forwarded.
  - Forward datagrams** from incoming to outgoing links.
    - The **data plane** processes and delivers packets based on the router/endpoint state, and constructs the forwarding table based on the routing algorithm.

## Principles

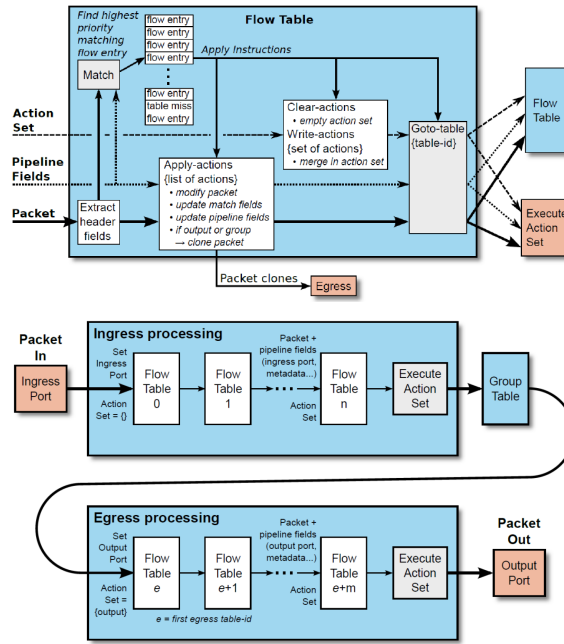
- Disaggregation:** separation of the data and control planes using an open interface.
  - Requires a **forwarding abstraction:** method to specify communications between the control and data planes.
    - E.g., **OpenFlow:** consists of *flow rules (Match-Action pairs)* which apply actions to packets with certain features.
  - Network operators able to **purchase** control & data planes from **different vendors**.
  - Shifting of control from vendors to operators** that can build networks to meet their users' needs.
  - Opportunities for fast innovation:** software control of the network can evolve independently of the hardware.
- Centralized control:** making the control plane *fully independent* of the data plane and *logically centralized*.
  - Centralized decisions are easier to make.
- Programmability:** enables performance optimizations and adaptability to protocol changes.
  - Data plane is implemented with a **forwarding pipeline:**
    - Each table focuses on a **subset of header fields** that might be involved in a given flow rule.
    - A packet is **processed by multiple tables** in sequence to determine how it is forwarded.
  - Easy network management:** programmers are able to debug/check behaviour more easily.
  - Rapid innovation:** each layer can evolve independently from one another.
  - Control shift** from vendor to operators and users.
- These principles use abstractions to modularize the network control problem.
  - Specification abstraction:** allow control apps to express the desired network behaviour without implementing it (*northbound interface*).
  - Distribution abstraction:** shield apps from distributed states, making control logically centralized (*middle layer*).
  - Forwarding abstraction:** allow any forwarding behaviour desired by apps while hiding details of the underlying hardware (*southbound interface*).

## Use cases

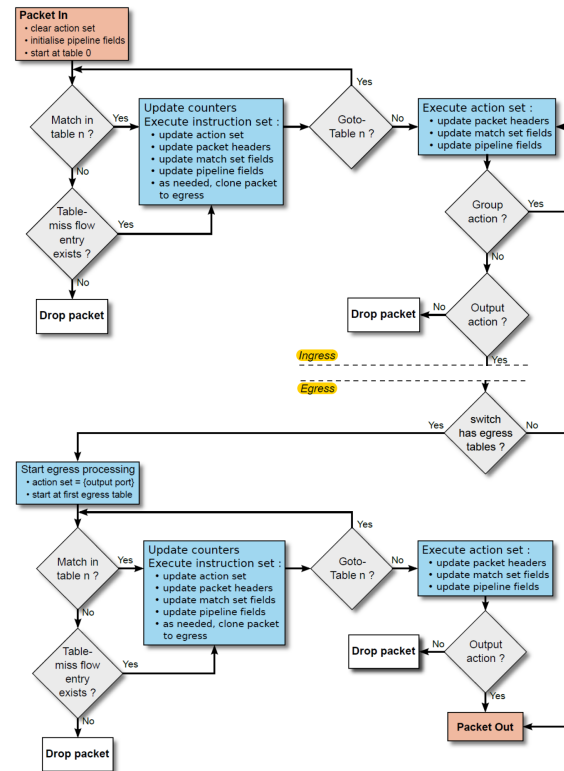
- Users of SDN:
  - Enterprises:** universities and private companies.
  - Cloud providers:** e.g., Google, Meta, Microsoft.
  - Network operators:** e.g., Comcast, AT&T, Singtel.
- Use cases of SDN:
  - Network virtualization:** enable networks to be programmatically created, managed, and torn down.
    - Single API entry point** to create, modify, and delete virtual networks.
    - Enables the virtual network to have its own private address space.
  - Switching fabrics:** enable lower costs and new features for cloud data centers.
  - Traffic engineering:** for wide-area links between data centers.
  - Software-defined WANs.**

## OpenFlow

- For each packet from a packet flow, it contains:
  - Header and header fields.**
  - Pipeline fields:** values attached to the packet during pipeline processing, e.g., ingress port and metadata.
  - Action:** an operation that acts on a packet, e.g., drop/forward/modify TTL.
  - Action set:** accumulated while processed by flow tables, and executed at the end of pipeline processing.
- A flow entry in a flow table contains:
  - Match field:** packets' header and pipeline fields are matched.
  - Priority:** used to choose from multiple matches.
  - Instruction set:**
    - Contains a list of actions to apply immediately.
    - Contains a set of actions to add to the action set (not applied immediately).
    - Enables modification of pipeline processing (e.g., going to another flow table, but backtracking to a previous table is not allowed).
  - Default entry:** flow entry for table-misses, by default packet will be sent to the control plane. If this is not present, the packet with a table-miss will be dropped.

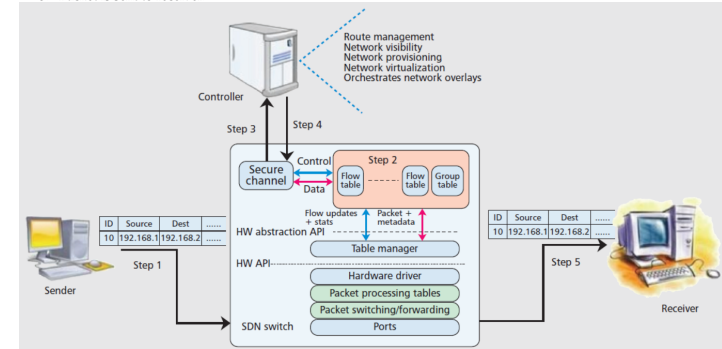


Packet flow through the processing pipeline



### Operation of SDN:

- Switch receives a packet from its input port.
- Switch tries to match packet against the flow entries in the flow table.
- If switch does not have any idea how to handle the packet, it forwards the packet to the controller.
- Controller writes new flow entry into the table.
- Packet is sent to receiver.



## Open network operating system (ONOS)

- Similar to other horizontally scalable cloud applications:
  - Consists of a set of **loosely coupled subsystems** (microservice architecture).
  - Contains a scalable and highly available **key-value store**.
- Architecture:**
  - Northbound interface (NBI):** for apps to stay informed about the network state and control the data plane.
  - Distributed core:** manages network state and notifies apps about relevant changes in state.
  - Southbound interface (SBI):** constructed from plugins, e.g., shared protocol libraries and device-specific drivers.
- Abstractions:**
  - Intent:** abstract network-wide, topology-independent programming constructs, where users can specify what they want but not how to achieve it.
  - Flow objective:** pipeline-independent; allows finer-grained control and device-centric programming constructs.
  - Flow rule:** pipeline-specific; supports both fixed-function and programmability.


## Interconnection

- Networks contain the following **basic units**:
  - Autonomous systems (AS):** network of interconnected routers, each identified by a unique AS number (ASN).
    - ASNs are assigned by the **Internet Assigned Numbers Authority (IANA)**, consisting of 5 **regional internet registries (RIRs)**.
    - They are controlled by a single administrative domain (but a company can have several ASNs).
  - Internet peering:** voluntary interconnection of administratively separate Internet networks, for exchanging traffic between the downstream users of each network.
    - Responsible for **bilateral interconnection**.
    - Depends on business relationships, which include:
      - Customer-provider relationships:**
        - Customer pays provider for access to the Internet, and reachability from anyone.
        - Provider provides **transit service** for the customer.
        - Most customers don't allow traffic through them (**nontransit AS**).
        - Some ASes provide transit between some providers (i.e., **selective transit**).
      - Multi-homing:** a customer with  $\geq 2$  providers.
      - Static routing** connects autonomous routing domains to the Internet  $\rightarrow$  customers do not need ASNs.
      - Peer-to-peer relationships** (i.e., *settlement-free peering*).
        - Peers provide transit between their respective customers.
        - Peers do not provide transit + often do not pay each other (i.e., **settlement-free relationship**).
        - Whether peering relationships are established depends on the following considerations:

To peer	Not to peer
Reduce upstream transit costs	Obtain and retain customers
Improve end-to-end performance	Working with your competition
Be the only way to connect customers to some part of the Internet	May require periodic renegotiation

- Internet exchange point (IXP):** an Ethernet switch where ASes freely interconnect to exchange traffic.
  - Responsible for **multilateral interconnection** (i.e., traffic exchange).
    - Ideally located at a **neutral, secure, accessible, safe, and expandable** place (i.e., **carrier-neutral**).
    - Each member (1) brings a router or (2) runs fibre/wireless from their data center to the IXP location, and connects it to the IXP switch.
  - Enable cheap peering by:
    - Saving upstream transit costs.
    - Keeping local traffic local.
    - Enable better network performance, quality of service, and scalability.
  - Networks are allowed to join the IXP if they have: (1) **address space**, (2) **ASN**, and (3) **transit agreements**.
  - Operation:** consortium model, representing all participants.
    - Costs** are agreed and **covered equally** by all participants.
    - Availability:** 24/7 cover provided by hosting location.

- Structure of the Internet: network of networks.



- Tier 1 commercial ISPs:** provide national and international coverage, e.g., Level 3, AT&T, NTT.
    - Have **access to the entire Internet**, only through its settlement-free peering links.
    - Peer with other tier 1s** to form a full mesh.
    - Have **no upstream provider**, they are the top of the customer-provider hierarchy.
  - Content providers:** private networks that connect their data centers to the Internet, often bypassing tier 1 and regional ISPs, e.g., Google.
  - Lower layer providers:** typically have national or regional coverage, and include a few thousand ASes.
    - Provide **transit to downstream customers**, but need at least one provider of their own.
  - Stub ASes:** do not provide transit service, merely connecting to upstream providers (e.g., consumers).
- The internet has a **valley-free property**, i.e., all valid AS paths are either:
- Single peak:** i.e., with a single tier-1 provider as root.
  - Single flat top:** i.e., with a peering at the root.
  - Any subpaths of the above.

## Inter-domain routing

- Challenges:**
  - Scale:** millions of routers and 200k+ prefixes.
  - Privacy:** ASes do not want to expose their internal topologies and business relationships with neighbours.
  - Policy:** no common notion of link cost + each AS needs control over its **traffic destination and peering**.
- Routing algorithms:**

Link state algorithm	Distance vector algorithm
Assumes all routers have knowledge of <i>complete topology and link cost</i>	Assumes all routers know <i>connected neighbours and their link costs</i>
Global or centralized	Decentralized
Employs <i>Dijkstra's algorithm</i>	Employs <i>Bellman-Ford algorithm</i> , iteratively exchanging information with their neighbours
E.g., <i>Open Shortest Path First (OSPF)</i>	E.g., <i>Routing Information Protocol (RIP)</i>

- Link state routing:**
    - (-) Nodes **divulge sensitive information**.
    - (-) **High bandwidth and storage overhead**, since topology information is flooded.
    - (-) **High processing overhead** in large networks, since the entire path is computed locally per node.
    - (-) Works only if **policies are shared and uniform**, since a standardized notion of *total distance* is required.
  - Distance vector routing:**
    - (+) **Hides details** of the network topology.
    - (-) Minimizes **some notion of total distance**.
    - (-) **Slow convergence**, certain parts of the network might have changed before the algorithm converges.
  - Path vector routing:** advertise the entire path, instead of just the distance metric, for each destination (e.g., BGP).
    - (+) **Fast loop detection:** nodes can detect a loop by checking if the path contains itself + discard paths with loops.
- Border Gateway Protocol (BGP):** the de facto inter-domain routing protocol.
- Allows subnets to **advertise their existence** to the rest of the Internet.
  - Allows ASes to **determine "good" routes** to other networks, based on reachability information and policy.
  - Provides **policy-based routing**.
    - Import policy:**
      - Filter unwanted routes** from neighbour.
      - Used to **rank customer routes over peer routes**.
      - Manipulate attributes to **influence path selection** (e.g., assign local preference to favoured routes).
    - Export policy:**
      - Filter out routes** you do not want to tell your neighbour.
      - Manipulate attributes to **control what neighbours see** (e.g., make paths look artificially longer).
    - Used by commercial ISPs to:
      - Fulfil bilateral agreements** with other ISPs: define who will provide transit for what.
      - Minimize monetary costs** / maximize revenue.
      - Ensure **good performance**.

### Types:

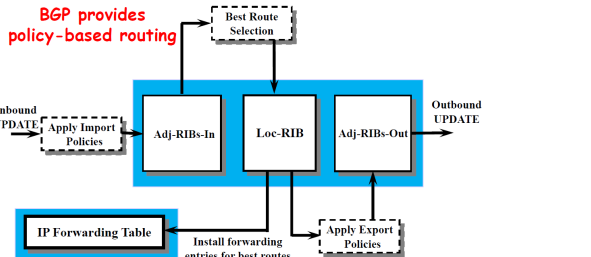
- Exterior BGP (eBGP) peering:** between BGP speakers in **different ASes**.
  - These routers should be **directly connected**.
  - When **ASX** advertises an IP prefix to **ASY**, **ASX promises** it will forward packets towards that prefix.
    - ASX can aggregate prefixes in its advertisement.
- Interior BGP (iBGP) peering:** between BGP peers **within an AS**.
  - Peers are **not required to be directly connected** (IGP e.g., RIP/OSPF, handles intra-AS peer connectivity).
  - iBGP peers must be **(logically) fully meshed**.
    - Used to **originate connected networks**.
    - Used to **pass on prefixes learned** from outside the AS (those from other iBGP peers are not passed on).

### BGP process:

- Establish BGP session.**
- Exchange all active routes.**
- Repeatedly exchange **incremental updates**.

### Getting entries into the forwarding table:

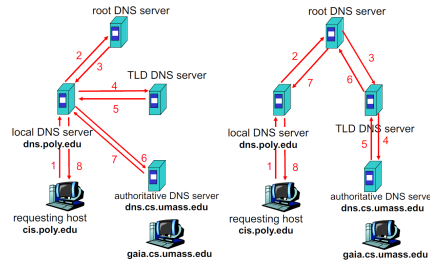
- Router becomes aware of prefix** via BGP route advertisements from other routers.
- Determine router output port** for prefix:
  - Choose the **best inter-AS route** from BGP (i.e., shortest AS-PATH).
    - BGP route = prefix + attributes = NLRI + path attributes.
  - Routing information bases (RIBs)** = Adj-RIBs-In + Loc-RIB + Adj-RIBs-Out: contains all routes.
    - Adj-RIBs-In:** unprocessed routes from peers via inbound UPDATE.
    - Loc-RIB:** selected local routes used by the router.
    - Adj-RIBs-Out:** selected for advertisement to peers.



- Use **OSPF** to find the best intra-AS route leading to the best inter-AS route.
  - Identify the **port number** for the best route.
- Add **prefix-port entry** to the forwarding table.
    - Hot potato routing:** if there is a tie, routers will try to differentiate these routes based on the distance that it needs to traverse within its network (by leveraging information in NEXT-HOP).
- Best route selection:**
    - Calculate **degree of preference**:
      - If the route is learned from an **internal peer**, use LOCAL\_PREF attribute or the pre-configured policy.
      - Else, use the pre-configured policy.
    - Route selection:** peers are recommended to use the highest **degree of LOCAL\_PREF**, and tie breaking on:
      - Smallest **number of ASNs in AS-PATH**.
      - Lowest **origin number in ORIGIN**.
      - Most preferred **MULTI\_EXIT\_DISC**.
        - For BGP peers to discriminate among entry points to a neighbouring AS → control inbound traffic.
      - Routes from **eBGP preferred over iBGP**.
      - Lowest **interior cost** based on NEXT\_HOP.
  - Messages:**
    - OPEN: opens TCP connection to peer and authenticates sender.
    - UPDATE: advertises new paths or withdraws old paths. Contains:
      - Withdrawn routes:** IP prefixes for the routes withdrawn.
      - Network layer reachability information (NLRI):** IP prefixes that can be reached from the advertised route.
    - KEEPALIVE: keeps connection alive in the absence of UPDATES; also ACKs OPEN request.
    - NOTIFICATION: reports errors in previous messages; also used to close connections.
  - Path attributes:**
    - Well-known mandatory:** e.g.,
      - ORIGIN: conveys the origin of the prefix.
      - AS-PATH: contains ASes through which NLRI has passed.
      - NEXT-HOP: indicates the IP address of the router in the next-hop AS.
    - Well-known discretionary.**
    - Optional transitive:** e.g.,
      - COMMUNITY: used to group destinations; each destination can be a member of multiple communities.
    - Optional non-transitive.**
      - Implementation rules:
        - Must recognize all **well-known** attributes.
        - Mandatory** attributes must be included in UPDATE messages that contain NLRI.
        - When BGP peers update **well-known** attributes, they must pass them to their peers.
  - BGP prefix hijacking:** re-routing internet traffic by falsely announcing ownership of IP addresses.
    - Blackhole:** traffic is discarded.
    - Snooping:** traffic is inspected and then redirected.
    - Impersonation:** traffic is sent to bogus destinations.
    - Subprefix hijacking:** traffic is redirected to a more specific prefix ( '.' traffic follows longest matching prefix).
- Prevention:
- Each AS **filters routes** announced by customers, **based on the prefixes that the customer owns**.
  - (-) Difficult to filter routes initiated from far away.
  - Route origin authorisation (ROA).**
  - Resource public key infrastructure (RPKI).**

## Peer-to-peer networks

- Centralized network models:**
  - Client/server model:** asymmetric, traditional communication model with ad-hoc clients and dedicated servers.
  - Delegation model:** first-hop server receives query as a delegation, forwards the query to other servers for processing, and eventually returns the result.
    - Query can be served either **iteratively** (left) or **recursively** (right).



- Iterative delegation is used by the **domain name system (DNS)**: prevents overloading the root DNS server.

### P2P network models:

- No always-on server/central entity.**
  - End systems directly communicate** with one another.
  - No a-priori knowledge** of the structure.
  - Flat architecture** / namespace.
- (-) Peers are **intermittently connected**, and change IP addresses → **unreliable service providers**.
- (-) Difficulty in **staying connected** and performing **resource lookup**.

- Applications:

### 1. File sharing:

- Napster:** based on a **central index server** which **knows all peers and files** in the network.
  - Delegation:** users register with the central server, which sends the list of files to be shared.
- Keyword-based searching:** server returns a list of details w.r.t. each file and the peer sharing it.
- (+) Consistent view of network.
- (+) Fast and efficient search.
- (+) Accuracy of search results in guaranteed.
- (-) Single point of failure + vulnerable to attacks.
- (-) Heavy computation required to handle queries.
- (-) Unreliable content, downloads from a single peer only.
- Gnutella:** pure p2p network.
  - An address of an active peer, obtained from an **out-of-band channel**, is required to join the network.
  - Once joined, peers learn about others and the topology of the network.
  - Queries are flooded** into the network, and peers **download directly** from others.
- (+) Fully decentralized.
- (+) Robust against node failures + less susceptible to denial of service attacks.
- (+) Open protocol, can be written on any OS.
- (-) Inefficient query flooding, wastes network and peer resources.
- (-) Inefficient network management.
- KaZaA:** consists of two kinds of nodes forming a 2-tier hierarchy, with top level = SNs and lower level = ONs.
  - Supernodes (SN):** peers with more resources/responsibilities.
    - Exchange information between themselves**, may change connections with other SNs in minutes.
    - Acts as a **hub for its children ONs**, keeping track of files in those ON peers.
    - Do not form a complete mesh**.
    - Do not cache information** from disconnected ONs.
  - Ordinary nodes (ON):** normal user/peer which sends requests through its parent SN.
    - Belongs to only 1 SN** at a time, and can change at will.
    - Not visible to other SN**.
    - Can be promoted to SN** if it has the desire and ability (sufficient resources: e.g., bandwidth, up time).

### 2. VoIP and messaging: e.g., Skype.


- Similar architecture to KaZaA: hierarchical overlay with SNs.
- Problem: **network address translation (NAT)** prevents an outside from initiating a call to inside peer.
  - Using Alice's and Bob's SNs, **relay** is chosen.
  - Each peer initiates session with relay.
  - Peers can now communicate through NAT via relay.
- Content distribution:** e.g., BitTorrent.
  - Builds a network (**swarm**) for *every file* that is being distributed.
- (+) Files can be sent to others as a link (i.e., .torrent file), which always refers to the same file.
  - For each file, one server (**seed**) initially hosts the original copy, before it is broken down into chunks.
  - Seed starts **tracker**: server keeping track of which seeds and peers are in the **swarm**.
  - Seed **creates torrent file** (with metadata: e.g., sizes + checksums of each chunk) and hosts it somewhere.
  - A new client **obtains the torrent file**.
  - Client **contacts tracker** and obtain the **torrent**: peers exchanging 256 kB chunks of the file.
  - Client **downloads/exchanges chunks** with peers.
    - Peers joining the torrent have no chunks, but will accumulate them over time.
    - Peers register with tracker to get a list of peers, and connects to a subset of them.
    - Peers may leave after they have obtained the entire file.
  - Pulling chunks:**
    - At any given time, peers have different subsets of file chunks.
    - Periodically, a peer may ask each neighbour for a list of chunks they have.
    - The peer sends requests for his/her missing chunks, and collects chunks via a **rarest first** strategy.
  - Pushing chunks:** tit-for-tat.
    - A peer sends chunks to 4 neighbours currently sending her chunks over at *the highest rate*.
    - The top 4 will be re-evaluated every 10 seconds.
    - Every 30s, the peer will randomly select another peer, and start sending chunks over to them.
- (-) **No searching possible:** no name service, but websites with "link collections" exist.
- Others: **Video streaming, p2p computation/storage, blockchains**.

### • Lookup services:

- Addressing:** supported by **structured networks/systems**, which allow for deterministic routing. Network structure determines where peers belong in the net, and where objects are stored.
- (+) Object location can be made efficient.
- (+) Each object is uniquely identifiable.
- (-) Need to know unique names.
- (-) Need to maintain structure required for addressing.
- Searching:** supported by **unstructured networks/systems**, where peers are typically free to join anytime and anywhere, and objects can be stored anywhere.
- (+) No need for unique names; more user friendly.
- (-) Hard to make efficient.
- (-) Need to compare actual objects, since query term may match multiple results.

### • Distributed hash tables (DHTs):

- Each node is responsible for  $\geq 1$  buckets.
  - Nodes communicate among themselves to find the responsible node.
  - Provide same abstraction supporting map operations, may differ in **namespace design and routing** in the overlay.
- Chord:**
    - Uses the **same hash function for both objects and nodes:** SHA-1.
    - Nodes are **organized in a ring**, keeping track of their predecessors and successors.
      - The **immediate successor** of each node is the closest node.
      - Each node is assigned an **identifier** in the range  $[0, 2^m - 1]$ .
    - Each node **maintains a finger table** with  $m$  shortcuts, with the  $i^{\text{th}}$  shortcut being  $\geq 2^{i-1}$  nodes away.
    - Node join:** finger table of other nodes are updated.
    - Routing:** shortcuts are taken based on the entries in the finger table.
    - Node leave:** each node periodically pings its 2 successors to see if they are still alive; if not, they will search for a replacement.



- Scalable content-addressable network (CAN):
  - Namespace is a  $d$ -dimensional torus, nodes only keep track of neighbours without storing shortcuts.
    - Node join:** suppose node  $B$  wants to insert a key-value pair  $(k, v)$ .
      - $B$  chooses a coordinate  $(a, b) = (h_i(k), h_j(k))$ .
      - The key-value pair  $(k, v)$  is routed to coordinate  $(a, b)$ .
      - If  $B$  is already a peer, the node who owns  $(a, b)$  stores  $x = a, y = b$ .
      - Else if the coordinate is owned by another node  $A$ ,  $A$ 's zone is split by half and  $B$  owns half of the cell. Splitting is done along the  $x$  dimension first, before the  $y$  dimension.
    - Routing:** routing is done by querying each node's 4 neighbours.
    - Node leave:**
      - The vacant zone is merged back to a neighbour, if possible.
      - Else, a neighbour node will temporarily handle multiple zones.
  - With higher dimensions (i.e.,  $d > 2$ ), routing table size and # hash functions  $\uparrow$ , but path lengths are shorter.

